

# reward\_noisy\_soft\_seed0 — preference validation loss (BCE)

